I built it to refuse.

Refusing was the risk.

If a reply looks like a fair housing matter, draft nothing.

Flag it, hand it to a person.

Broken fan: drafted reply in two seconds.

Grab bar for a resident’s mother: nothing. Wait for a human.

Two residents. Two speeds. The difference being that one of them mentioned a disability.

Frictionless service for ordinary requests, a slower queue for protected class requests.

Analysts describe that as possible disparate impact.

Algorithmic tools fall under the Fair Housing Act. Disparate outcomes are actionable regardless of intent.

Not knowing is not a *defense.*

A language model ~~cannot reliably tell~~ a pet rent question from an assistance animal request.

So detection runs on fixed rules, before the model.

Every category drafts at the same speed.

The flag changes what the manager sees.
It ~~does not change~~ what the resident waits for.

Audit where your automation does less, not only where it does more.

Refusing is a behavior. It is ~~not automatically~~ the safe one.

Found by someone asking what a button did.

Not by a test suite, and not by me.
Not legal advice. Ask your own counsel.

SAMBHAV JAIN